

Asynchronous Federated and Reinforcement Learning for Mobility-Aware Edge Caching in IoVs

Journal:	<i>IEEE Internet of Things Journal</i>
Manuscript ID	IoT-32000-2023
Manuscript Type:	Regular Article
Date Submitted by the Author:	29-Aug-2023
Complete List of Authors:	Jiang, Kai; Wuhan University Cao, Yue; Wuhan University, Song, Yujie; Wuhan University Zhou, Huan; China Three Gorges University, Wan, Shaohua; University of Electronic Science and Technology of China, Shenzhen Institute for Advanced Study Zhang, Xu; Xi'an University of Technology, Computer Science and Technology
Keywords:	Intelligent Transportation Systems < Sub-Area 3: Services, Applications, and Other Topics for IoT, Cyber-Physical Systems < Sub-Area 3: Services, Applications, and Other Topics for IoT, Vehicular Networks < Sub-Area 2: Communications and Networking for IoT, Smart Cities < Sub-Area 3: Services, Applications, and Other Topics for IoT, Service Functions and Management < Sub-Area 3: Services, Applications, and Other Topics for IoT, Mobile and Ubiquitous Systems < Sub-Area 3: Services, Applications, and Other Topics for IoT

Asynchronous Federated and Reinforcement Learning for Mobility-Aware Edge Caching in IoVs

Kai Jiang, Yue Cao, *Senior Member, IEEE*, Yujie Song, Huan Zhou, *Member, IEEE*,
Shaohua Wan, *Senior Member, IEEE*, and Xu Zhang *Member, IEEE*

Abstract—Edge caching is a promising technology to reduce backhaul strain and content access delay in Internet-of-Vehicles (IoVs). It pre-caches frequently-used contents close to vehicles through intermediate roadside units. Previous edge caching works often assume that content popularity is known in advance or obeys simplified models. However, such assumptions are unrealistic, as content popularity varies with uncertain spatial-temporal traffic demands in IoVs. Federated learning (FL) enables vehicles to predict popular content with distributed training. It preserves the training data remain local, thereby addressing privacy concerns and communication resource shortages. This paper investigates a mobility-aware edge caching strategy by exploiting asynchronous FL and Deep Reinforcement Learning (DRL). We first implement a novel asynchronous FL framework for local updates and global aggregation of Stacked AutoEncoder (SAE) models. Then, utilizing the latent features extracted by the trained SAE model, we adopt a hybrid filtering model for predicting and recommending popular content. Furthermore, we explore intelligent caching decisions after content prediction. Based on the formulated Markov Decision Process (MDP) problem, we propose a DRL-based solution, and adopt neural network-based parameter approximations for the curse of dimensionality in RL. Extensive simulations are conducted based on real-world data trajectory. Especially, our proposed method outperforms FedAvg, LRU, and NoDRL, and the edge hit rate is improved by roughly 6%, 21%, and 15%, respectively, when the cache capacity reaches 350 MB.

Index Terms—Edge caching; federated learning; content prediction; stacked autoencoder; deep reinforcement learning.

I. INTRODUCTION

With the popularity of vehicles on the road and the advancements in V2X technologies, the transportation system has evolved into a data-driven intelligent era [1]. Internet-of-Vehicles (IoVs) enable vehicles with diverse vehicular services, from multimedia to safety-related applications. Technically, these attractive applications demand access to vast Internet contents, relying heavily on ultra-reliable and low-latency communications for content delivery. The significant surge in delay, unreliability, and redundant traffic have rendered cloud-

based processing architectures infeasible due to the limited backhaul capacity and long transmission distance [2]–[4].

Fortunately, large-scale analysis has revealed that different contents often necessitate different priorities, which has prompted the emergence of edge caching in IoVs [5]. Expressly, edge caching is a promising technology to reduce backhaul strain and content access delay. It pre-caches frequently-used contents close to vehicles by sinking cloud functions to intermediate Roadside Units (RSUs). Hence, vehicles can directly fetch content from caching-enabled RSUs, instead of duplicate transmissions from remote cloud servers.

Generally, the finite caching capacity of an RSU makes the system performance heavily relies on well-designed caching strategies. The key to strategy design lies in knowing content popularity, which reflects content access preferences in the content library. Previous works [6], [7] often assume that content popularity is known in advance or obeys simplified models, like *Zipf* distribution and its variants. Nevertheless, such assumptions are unrealistic in practice, as content popularity exhibits non-stationary with uncertain spatial-temporal traffic demands in IoVs. Therefore, to facilitate edge cache utilization, predicting content popularity is urgently required.

Recent advancements in data analysis and computing power have provided opportunities to improve predictions from a data-driven perspective. Based on this, Machine Learning (ML) has exhibited great potential in extracting dynamic features of content popularity by training Vehicular User (VU) data [8]. However, despite the continuous progress achieved in ML, adopting ML for content prediction in IoVs still faces two issues: 1) The regional content popularity and validity are constantly changing. As vehicles connected to an RSU move and traverse the cell coverage rapidly, cached content can easily be outdated. An effective prediction method should be context- and mobility-aware to improve cache performance. 2) Most ML methods operate on the centralized framework, where vehicles must transmit VU data to an RSU for training. It is inconsistent with the growing privacy concerns and limited communication resources in data transmission. The VU data usually involves sensitive information, and VUs are reluctant to share their data with others directly [9], [10].

To address these issues, Federated Learning (FL) has attracted widespread concern as a privacy-preserving solution for distributed learning. Coordinated by RSUs, FL enables a shared global model update by aggregating local model parameters from vehicles. Each vehicle conducts training on its local VU data and uploads only model parameters to the RSU in proximity, instead of raw data [11], [12]. However,

K. Jiang, Y. Cao (corresponding author), and Y. Song are with the School of Cyber Science and Engineering, Wuhan University, Wuhan 430000, China. (e-mail: kai.jiang, yue.cao, Y. Song@whu.edu.cn).

H. Zhou is with the School of Computer Science, Northwestern Polytechnical University, Xi'an 710129, China. (e-mail: zhouhuan117@gmail.com).

S. Wan is with the Shenzhen Institute for Advanced Study, University of Electronic Science and Technology of China, Shenzhen 518110, China. (e-mail: shaohua.wan@uestc.edu.cn).

X. Zhang is with the Department of Computer Science and Engineering, Xi'an University of Technology, Xi'an 710048, China (e-mail: zhangxu@xaut.edu.cn).

traditional FL operates aggregation synchronously, this leads to resource-wasting as the RSU must wait for all local updates from vehicles before aggregation. The synchronous manner raises two other issues. First, due to vehicle heterogeneity and communication uncertainties, some lagging vehicles may experience delays in local training, becoming "stragglers" [13]. Second, vehicle mobility will hinder model aggregation, as vehicles might leave the cell coverage of RSU before uploading their local updates. Such unavailability of vehicle inevitably disrupts the timely contribution of certain vehicles to the global model. As a remedy, exploring asynchronous FL becomes imperative to enhance content prediction adaptability, mitigate the impact of stragglers and ensure up-to-date global models.

In addition, although the prediction results offer valuable insights for caching decisions, the diversity of content often results in predicted popular content far surpassing the finite cache capacity of an RSU. To facilitate edge caching utilization and maximize the benefits of content prediction, intelligent content replacement mechanisms are indispensable in edge caching [14]. Therefore, the RSU must determine where to cache the predicted popular contents and whether to replace already cached contents. However, complex content requests and vehicle mobility pose new challenges for caching decision-making. Deep Reinforcement Learning (DRL) effectively addresses these challenges. With solid cognitive ability, DRL handles dynamic control problems with high-dimensional and time-varying features. It is especially suited for sequential decision-making with long-term objectives under uncertainty [15], [16]. Thus, we can execute content decision-making in IoVs by exploiting DRL.

Based on the above considerations, this paper investigates the mobility-aware edge caching by asynchronous FL and DRL. The main contributions are summarized as follows:

- 1) We conduct popular content prediction by Stacked AutoEncoder (SAE) models among vehicles and the RSU. The local updates and global aggregation of SAE models follow an asynchronous FL framework. Unlike traditional synchronous FL awaiting all local updates, this framework employs asynchronous aggregation, where the central RSU updates its global model immediately for each arrived local update. Meanwhile, we cope with the incurred stragglers in training and improve model convergence by several interventions.
- 2) Utilizing the latent features extracted by the SAE model, we adopt a hybrid filtering model for predicting popular content. Based on the content ratings and personal information, this hybrid filtering model combines content-based collaborative filtering and demographic information during runtime. It determines the distance between two contents or two VUs through similarity metrics.
- 3) Furthermore, to maximize the benefits of content prediction, we formulate the caching process in an RSU as an Markov Decision Process (MDP) problem, intending to minimize the long-term content delivery delay in the system. Our proposed solution is based on DRL with adaptive and foresighted considerations, adopting neural network-based parameter approximations to circumvent the curse of dimensionality in RL.

The remainder of this paper is organized as follows. Section II and Section III describe the related work and the system model, respectively. Section IV proposes a popular content prediction method under the FL framework. Section V explores caching decisions based on DRL to maximize the benefits of content prediction. Moreover, simulation results are analyzed in Section VI. Conclusion is summarized in Section VII.

II. RELATED WORK

A. Distributed Optimization

With the popularity of vehicles on roads and their enhanced computing capabilities, training models directly on these distributed vehicles has drawn increasing attention. Thereinto, conventional multi-task models [17] are unsuitable for edge vehicle training due to the assumption that all vehicles partake in each training round. This assumption is unrealistic as vehicles may frequently go offline during training for network unreliability or other factors. In light of this, distributed optimization is a critical paradigm to address the complexities of large-scale model training across vehicular networks. It aims to collectively optimize a global objective function while each vehicle performs local computations based on its own data. These local computations are then aggregated to achieve a global solution.

Nevertheless, in recent years, various distributed optimization methods [14], [18]–[20] encountered slow convergence, inefficiency, and straggler issues when deployed to real-world scenarios that involve large-scale and highly heterogeneous data with rich features. For instance, Zhou et al. [14] concentrated on optimizing cache replacement and multipoint collaboration by distributed multi-agent RL. Notably, however, this method exhibits exponential space complexity with respect to the number of agents. The computation of the Nash equilibrium significantly dominates the training time. Although they introduced parameter approximation to mitigate the convergence complexity, the convergence pattern remains nonlinear with the number of agents. Jin et al. [19] devised a distributed framework with alternating local and central learning by the soft confidence-weighted classifier. However, this asynchronous method assumes normally distributed local data, limiting non-convex neural network objectives. Meanwhile, it also lacks theoretical convergence assurances and necessitates transmitting a fraction of local data from each device to the centralized server.

B. FL for Popular Content Prediction

FL, an extension of distributed optimization, offers a promising solution to tackle device diversity in ad-hoc networks. In contrast to distributed optimization, FL emphasizes privacy preservation and enables training on sensitive data. Meanwhile, FL is tailored to handle challenges like non-uniform data distribution, communication bottlenecks, and network unreliability.

Recently, FL has acquired noteworthy application significance in edge caching, with several promising endeavours focused on popular content prediction [11], [12], [21]–[24]. Specifically, Yu et al. [11] introduced a mobility-aware

proactive edge caching strategy using FL, where a context-aware adversarial AE is employed for content popularity prediction. This work dramatically enhances cache efficiency and eliminates the need for centralized data processing, but coordinating vehicle learning through synchronous aggregation is unsuitable for highly dynamic vehicular edge network scenarios. Xie et al. [21] introduced an asynchronous aggregation for federated optimization involving weighted averaging of the global model. However, they did not consider the practical conditions where edge devices must deal with continuous streaming data. Furthermore, Wang et al. [23] focused on employing FL in edge caching to improve data security with consider the potential private data leakage in the trained model. They incorporated gradient clipping and model parameter limitation in model training, and proposed a privacy-preserving method for content popularity prediction by FL and Wasserstein generative adversarial network. Jiang et al. [24] introduced an FL-based model integration approach that utilizes user context information for efficient user clustering. They employed the distributed approximate Newton algorithm with stochastic variance reduced gradient to learn the global popularity prediction model based on local models. However, this method does not consider user mobility and potential straggler issues during FL training.

C. DRL for Edge Caching Decision-Making

With the revival of artificial intelligence, DRL has been widely used in caching decision-making [5], [7], [15], [25]–[27]. These methods enable the acquisition of key attributes, including user historical request information and content popularity.

Specifically, Jiang et al. [5] reviewed DRL-driven edge caching comprehensively and discussed its implementation and outlook. Considering the vehicle mobility and stringent deadline constraints, Tan and Hu [15] embarked on a bold step in designing, analyzing and optimizing the cooperative coded caching placement and computing allocation at both the vehicle and RSU levels in vehicular networks. They configured edge caching problems with a DRL-based multi-timescale framework, and developed mobility-aware reward estimation in a large timescale model to relieve spatial complexity. Similarly, Qiao et al. [25] harnessed a comparable framework to optimize the content placement and delivery in the vehicular edge networks, with the aid of flexible trilateral cooperations among a macro-cell station, RSUs, and vehicles. Moreover, in [26], Chen et al. considered the service caching and task dependency, and addressed the collaborative service caching problem in a digital twin-empowered edge architecture by asynchronous advantage actor-critic. In [27], Kirilin et al. exploited cache admission, and presented a model-free DRL-driven method for determining the admission of requested objects into the CDN cache. In contrast to prior methods that rely on a limited set of criteria, this method weights a large set of features, including the object size, recency, and frequency of access. The aforementioned works highlight the efficacy of DRL for decision-making. Nevertheless, their reliance on centralized ways may not always be practical in real-world

TABLE I
NOTATIONS AND SYMBOLS

Notation	Explanation
\mathcal{N}	The set of all RSUs
\mathcal{F}	The index set of available contents
s_f	The size of each content f
G_i	The limited cache capacity of RSU i
$T_{u,i}$	The transmission delay between vehicle u and RSU i
$T_{i,j}, T_{i,N+1}$	The transmission delay between RSU-RSU and cloud server-RSU
$r_{u,i}$	The wireless downlink data rate between vehicle u and RSU i
$\mathcal{T}_{f,i,j}$	The total delay through different links
$a_{f,i,j}^t$	The caching decision of any RSU i for the requested content f
ξ	The Poisson distribution parameter
V_u	The velocity of vehicle u
μ, σ	The mean and standard deviation of vehicle velocities
P_u^t	The position of vehicle u

scenarios due to distributed network topology and privacy concerns.

D. Our Motivation

Inspired by existing studies, it's evident that fewer works have considered both vehicle mobility and privacy concerns when designing adaptive and foresighted caching decision-making. This motivates us to explore a mobility-aware edge caching strategy by exploiting asynchronous FL and DRL. We implement a novel asynchronous FL framework, and conduct popular content prediction by SAE models among vehicles and the RSU. Compared to existing synchronous FL works, our work employs asynchronous aggregation while coping with straggler issues through several interventions. In addition, to maximize the benefits of content prediction and overall edge cache utilization, our work proposes a DRL-based solution based on the formulated MDP problem. It adopts neural network-based parameter approximations to circumvent the curse of dimensionality in RL. Finally, based on real-world trajectory, we conduct extensive simulations to verify the effectiveness of the proposed method.

III. SYSTEM MODEL

This section presents the edge caching architecture, including network infrastructure, content delivery, and vehicle mobility models. The main notations are listed in **Table I**.

A. Network Architecture

We consider an IoV architecture in the highway scenario comprising a cloud server and several small cells, each of which is equipped with a caching-enabled RSU, as illustrated in **Fig. 1**. All these RSUs in a cluster form are deployed equidistantly along the one-way street, and the central RSU is




Fig. 1. Edge Caching-empowered IoV Architecture.

responsible for subsequent FL training. The set of all RSUs is denoted by $\mathcal{N} = \{1, 2, \dots, N\}$. Besides, let $\mathcal{F} = \{1, 2, \dots, F\}$ represent the total available contents, and s_f is the size of each content. The cloud server (denoted by $N + 1$) has abundant cache capacity to cache all contents. Each RSU $i (i \in \mathcal{N})$ is limited to a cache capacity of G_i , which indicates the maximum number of caching contents. Vehicles travel in the single direction along the road and frequently request content via V2X. These vehicles are assumed to be located in at least one cell coverage and will be served by the associated RSU therein. Particularly, each vehicle possesses its unique local dataset, which includes access requests and rating information for contents, as well as context information of VUs (age, gender, *etc.*). Meanwhile, we assume the system operates in multiple time slots (*e.g.*, batching or round) $t \in \{1, \dots, \Gamma\}$. Each time slot allows a vehicle to request a content, and the caching decision of each RSU is updated periodically. Notably, vehicle number in a cell coverage changes in each slot due to the mobility.

Each RSU caches a small portion of contents to eliminate the redundant traffic, ensuring that content requests can be accommodated nearby. Particularly, RSUs can determine where content requests are answered and which local caches should be replaced. Once a vehicle sends an access request within the cell coverage, the local RSU checks its cache space to determine whether the desired content is already cached. If not, the local RSU can retrieve it from adjacent RSUs or directly download it from cloud servers, and deliver it to the vehicle later. Additionally, all RSUs can replace their obsolete content with the popular one in each time slot. This involves a sequential decision-making problem, including efficient decision optimization under constraints.

B. Content Delivery Model

Transmission delay of content f from any requested vehicle u to its connected RSU i is calculated as $T_{u,i} = s_f / r_{u,i}$, where $r_{u,i}$ is the wireless downlink data rate between vehicle u and RSU i . Here, multiple channels are considered within each RSU, all with the same bandwidth allocation [14]. Considering the large-scale fading, $r_{u,i}$ is derived as follows:

$$r_{u,i} = B_{u,i} \log_2 \left(1 + \frac{P_i g_{u,i}}{\sigma^2 + \sum_{v \in \mathcal{U} \setminus \{u\}} P_i g_{v,i}} \right), \quad (1)$$

where $B_{u,i}$ denotes the channel bandwidth, P_i denotes the transmission power of RSU i , σ^2 and $g_{u,i}$ are the noise power and the channel gain, respectively. Notably, interference management is considered for this wireless transmission, and we neglect the transmission delay of sending request identifiers as it involves tiny data size and occurs over high link rates in most instances. Furthermore, as cloud server-RSU and RSU-RSU links are wired optical cables, the corresponding transmission delay $T_{i,N+1}$ and $T_{i,j}$ are set to constant values.

Here, we introduce binary variable $a_{f,i,j}^t \in \{0, 1\}$, $j \in \mathcal{N} \cup \{N + 1\}$ to denote the decision of any RSU i for the requested content f in time slot t . Upon a content request, the content delivery delay is analyzed as follows:

- $a_{f,i,i}^t = 1$ indicates that content f is cached in the RSU i at slot t and can be delivered to the vehicle directly. In this case, the total delay $\mathcal{T}_{f,i,i}$ is just the transmission delay $T_{u,i}$ from the RSU to the vehicle;
- $a_{f,i,j}^t = 1$ ($j \in \mathcal{N} \setminus \{i\}$) indicates that content f is not cached in the RSU i , but at least one of RSUs in the cluster has cached this content, thus the local RSU can inquire and obtain it from the neighbouring RSU j . Then, the total delay $\mathcal{T}_{f,i,j}$ consists of the transmission delay $T_{i,j}$ from the neighbouring RSU j to the local RSU i and the transmission delay $T_{u,i}$;
- $a_{f,i,N+1}^t = 1$ indicates that there is no expected content in all RSUs, and the RSU i has to forward the request identifier to the cloud server for processing at slot t . In this way, the total delay $\mathcal{T}_{f,i,N+1}$ consists of the transmission delay $T_{i,N+1}$ from cloud server to RSU i and the transmission delay $T_{u,i}$.

C. Vehicle Mobility Model

Vehicle mobility is an intrinsic feature in IoVs-related research. For each time slot t , we assume that the number of arrived vehicles in a cell coverage follows the Poisson distribution with parameter ξ . Here, parameter ξ is the average arrival frequency, which reflects the mobility intensity. Notably, due to the tiny slot size, we do not consider switching among cell coverages, and vehicle characteristics (velocity, position) remain unchanged throughout a time slot. This ensures a consistent number of arrived vehicles for each cell coverage in a time slot. Particularly, in the free-flow traffic state, vehicle velocities are assumed to follow independent and identically distributed (i.i.d.) distributions, as drivers can regulate them autonomously. Therefore, we hypothesize that the velocity of vehicle u is derived from a Gaussian distribution, with the probability density function as:

$$f(V_u) = \frac{1}{\sigma \sqrt{2\pi}} e^{-\frac{(V_u - \mu)^2}{2\sigma^2}} \quad (2)$$

where V_u is the vehicle velocity; μ and σ is the mean and standard deviation of vehicle velocities, respectively.

Without loss of generality, we assume that vehicle velocities in urban districts are restricted within the upper and lower bounds (*i.e.*, $V_{\min} < V_u < V_{\max}$ for vehicle u). To avoid

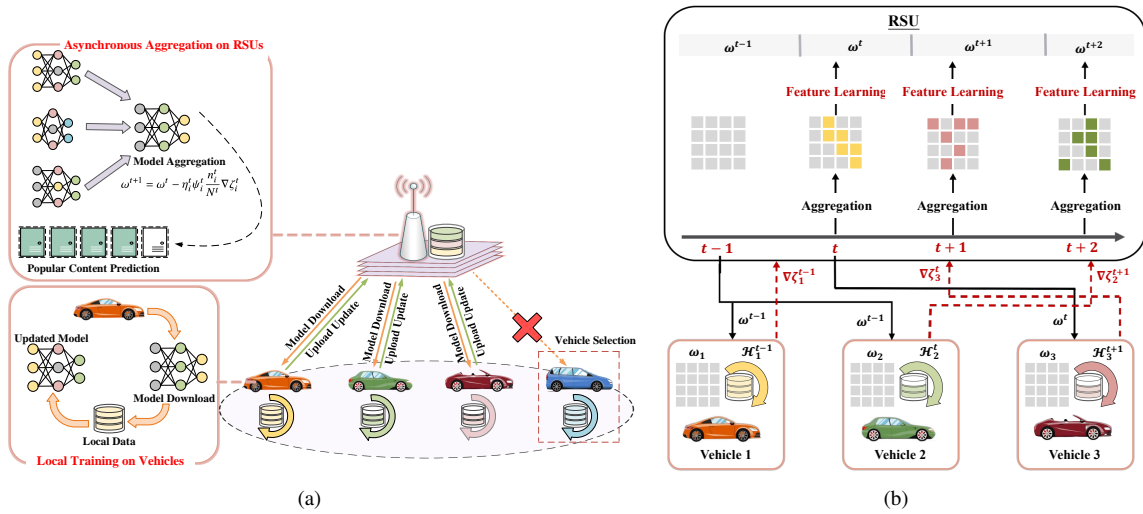


Fig. 2. (a) Mobility-Aware Asynchronous FL based Model Training Process. (b) Asynchronous Aggregation with the Representation Learning.

negative or near-zero velocities, we shall truncate the above Gaussian distribution as:

$$\hat{f}(V_u) = \begin{cases} \frac{f(V_u)}{\int_{V_{\min}}^{V_{\max}} f(s) ds} = \frac{2f(V_u)}{\text{erf}\left(\frac{V_{\max}-\mu}{\sigma\sqrt{2}}\right) - \text{erf}\left(\frac{V_{\min}-\mu}{\sigma\sqrt{2}}\right)}, & V_{\min} < V_u < V_{\max} \\ 0, & \text{otherwise} \end{cases} \quad (3)$$

where $\text{erf}(x) = \left(\frac{2}{\sqrt{\pi}}\right) \int_0^x e^{-t^2} dt$ is the Gaussian error function under the mean μ . Compared to the traditional Gaussian distribution or fixed velocity setting, the truncated version is more feasible and aligns better with practical traffic conditions.

Furthermore, we assume that vehicles maintain continuous connections with the RSU within its cell coverage. The connection time can be rather longer if vehicles travel at lower velocities. Here, let P_u^t represent the position of vehicle u in time slot t (i.e., the traversed distance within the coverage of its connected RSU). Hence, for each vehicle u , the connection time in the cell coverage of the current RSU is expressed as:

$$T_{u,conn}^t = (F_d - P_u^t) / V_u^t \quad (4)$$

where F_d is the diameter of cell coverage.

IV. MODEL TRAINING AND CONTENT PREDICTION

This section proposes a popular content prediction method under the FL framework. We employ asynchronous FL for local updates and global aggregation of SAE models. The trained SAE model extracts latent features from VUs and contents. These features are used to construct a hybrid filtering model for recommending popular content.

A. Mobility-Aware Asynchronous FL based Model Training

FL enables collaborative training of SAE models among vehicles and the RSU, and preserves vehicle data remain local. The schematic overview of model training is illustrated in Fig. 2(a). The central RSU distributes the global model (i.e., SAE model) to selected vehicles at the beginning of each round. Each vehicle initializes its local model with the received global

model and trains it iteratively by local data. Following the local training, vehicles upload their local updates to the central RSU, which then performs asynchronous aggregation to update the global model. Notably, we cope with the incurred stragglers in FL and improve model convergence by several interventions. Along with considering vehicle mobility, we incorporate regularized local loss functions and dynamic learning rates during local training. Additionally, we balance the previous and the current local gradients using a decay coefficient, and implement representation learning to mitigate the performance impact of asynchronous optimization.

1) *Vehicle Selection and Model Download*: The high-mobility may impose vehicles traversing the cell coverage rapidly, with insufficient time to complete the FL training. Thus, the central RSU should select eligible vehicles for FL training at the start of each round. The selection process considers favorable channel conditions, sufficient local training data, and the remaining connection time within the cell coverage. Thereinto, connection time determines whether a vehicle can participate in FL training and transmit its local update to the central RSU. As displayed in Eq. (4), the connection time of a vehicle largely depends on its position and velocity. Let T_{tra} and T_{inf} be the average training time and inference time in each slot, respectively. For each vehicle u , only if its connection time within the cell coverage exceeds the sum of these two means (i.e., $T_{u,conn} > T_{tra} + T_{inf}$), will it have the opportunity to participate in FL training in the current round. Hereafter, the RSU distributes the global model to selected vehicles, enabling them to improve it through local training.

2) *Training on Vehicles*: For round t , a portion of vehicles is selected to download the global model parameter ω^t from the central RSU. Then, vehicles perform gradient optimizations and refresh their local model parallel based on the local data for multiple iterations. Let \mathcal{I}^t denote the set of selected vehicles, and the local data captured on each vehicle u is \mathcal{H}_u^t , with $h_u^t = |\mathcal{H}_u^t|$ being the number of data samples.

The empirical local loss function of vehicle u is defined as:

$$\ell_u(\omega_u^t) = \frac{1}{h_u^t} \sum_{x \in \mathcal{H}_u^t} (x - \hat{x})^2 = \frac{1}{h_u^t} \sum_{x \in \mathcal{H}_u^t} f(\omega_u^t; x) \quad (5)$$

where $f(\omega_u^t; x)$ is the loss function for data point x under local model ω_u^t , and \hat{x} is the input reconstructed by SAE.

Notably, a regularization operation is implemented to ensure algorithm convergence during the local training. The regularized local loss function is defined as:

$$s_i(\omega_u^t) = \ell_u(\omega_u^t) + \frac{\lambda}{2} \|\omega_u^t - \omega^t\|^2 \quad (6)$$

where the second term is the regularization term to avoid the overfitting, with λ is the regularization parameter.

Generally, lagging vehicles (*i.e.*, stragglers) are inevitable in FL training due to statistical heterogeneity and communication uncertainty. These stragglers may lag in uploading their updates as they have more data or poor communication conditions in the previous round. We define the gradient of the regularized local loss function as "local gradient", and distinguish the local gradient of stragglers in the previous round as "delayed local gradient". Empirical studies have revealed that stragglers negatively affect the accuracy and convergence of the global model. This is attributed to the parameter differences used in calculating the current local gradient compared to the delayed local gradient. To this end, balancing the previous and current local gradients during the aggregation in the current round is essential. That is,

$$\nabla \zeta_u^t = \nabla s(\omega_u^t) + \delta \nabla s_u^{(pre)}. \quad (7)$$

where $\nabla \zeta_u^t$ is the aggregated local gradient of vehicle u , δ is the decay coefficient, $\nabla s(\omega_u^t)$ and $\nabla s_u^{(pre)}$ are local gradient and delayed local gradient, respectively. Notably, $\nabla s_u^{(pre)} \rightarrow 0$ for uploading updates duly in the previous round.

Furthermore, to cope with the stragglers in asynchronous optimization, we incorporate a dynamic local learning rate into the local update. The rationale is that the corresponding learning step size for stragglers should be increased. With η_u^t being the time-related dynamic local learning rate of vehicle u in round t , the closed-form iterative solution for the local model update is defined by:

$$\omega_u^{t+1} = \omega_u^t - \eta_u^t \nabla \zeta_u^t \quad (8)$$

where multiplier $\eta_u^t = \eta_d \max\{1, \log(\bar{d}_u^t)\}$, η_d is an initial learning rate set for all vehicles, and $\bar{d}_u^t = \frac{1}{t} \sum_{\tau=1}^t d_u^\tau$ is the average time cost of the past t rounds. Following that, the actual learning step size is adjusted based on previous delays. By assigning larger step sizes to stragglers, this adaptive scaling effectively mitigates the adverse impact of stragglers on model convergence.

The local iteration process persists until the maximum number of epochs in FL training. After that, vehicle u uploads its local update (*i.e.*, $\eta_u^t \nabla \zeta_u^t$) to the RSU. The global loss function is formally defined as:

$$\mathcal{L}(\omega^t) = \sum_{u \in \mathcal{I}^t} \frac{n_u^t}{N^t} \ell_u(\omega_u^t) \quad (9)$$

where $N^t = \sum_{u \in \mathcal{I}^t} |\mathcal{H}_u^t|$ is the number of whole data among selected vehicles, $\frac{n_u^t}{N^t}$ is the proportion of local data to all data. Thus, vehicles with more data contribute more to the aggregated parameters, thereby preventing model contamination by vehicles with less data. Our goal of FL is to obtain a global model ω^* with:

$$\omega^* = \arg \min_{\omega^t} \mathcal{L}(\omega^t) \quad (10)$$

3) *Asynchronous Aggregation on the Central RSU*: The global model in central RSU gets updated by aggregating the local update from vehicles in each round. Traditional FL operates synchronously, which leads to resource-wasting as the RSU must wait for all local updates before aggregation. In contrast, we perform asynchronous FL where the central RSU updates its global model immediately for each arrived local update, without waiting for other vehicles to upload. Specifically, for the local update from vehicle u , the global model will be updated by:

$$\begin{aligned} \omega^{t+1} &= \omega^t - \frac{n_u^t}{N^t} (\omega_u^t - \omega_u^{t+1}) \\ &= \omega^t - \frac{n_u^t}{N^t} (\omega_u^t - (\omega_u^t - \eta_u^t \nabla \zeta_u^t)) \\ &= \omega^t - \eta_u^t \frac{n_u^t}{N^t} \nabla \zeta_u^t \end{aligned} \quad (11)$$

Meanwhile, inspired by [11], we recognize the significance of incorporating different weights during model aggregation, with considering vehicle mobility. Consequently, to enhance the global model accuracy and reduce the content transmission delay, the aggregation weight ψ is introduced with the intuition that vehicles with longer available training time should contribute more and be endowed with larger weights. Based on the wireless downlink rate and remaining traversed distance in the cell coverage, the normalized aggregate weight of vehicle u can be derived as:

$$\psi_u^t = \alpha \frac{(F_d - P_u^t)}{F_d} + (1 - \alpha) \frac{r_{u,i}^t}{\max_{u \in \mathcal{I}^t} (r_{u,i}^t)} \quad (12)$$

where α is the tradeoff factor between wireless downlink rate and remaining traversed distance. As a result, the aggregated global model in Eq. (11) can be rewritten by incorporating the aggregate weight with it:

$$\omega^{t+1} = \omega^t - \eta_u^t \psi_u^t \frac{n_u^t}{N^t} \nabla \zeta_u^t \quad (13)$$

In addition, to mitigate the performance impact of asynchronous optimization, we implement representation learning on the central RSU. This enables us to capture embedding representations for most features, which are learned simultaneously with other parameters in the SAE model. The central RSU executes aggregation immediately for each arrived local update, and performs representation learning on aggregated parameters to derive a cross-vehicle feature representation. The illustration of asynchronous aggregation with representation learning is displayed in **Fig. 2(b)**. Specifically, our approach draws inspiration from attention mechanisms for feature identification and representation. Meanwhile, we in-

Algorithm 1: Mobility-Aware Popular Content Prediction
 based on Asynchronous FL

1 Central RSU Execution:

Input: tradeoff factor α , the position and velocity of vehicles, wireless downlink rate for vehicles.

2 Initialize $\omega \leftarrow \omega^0$
3 for each round t **do**

4 I^t : the set of selected vehicles

5 for each vehicle u **in parallel do**

6 $T_{u,conn}^t = (F_d - P_u^t) / v_u^t$

7 **if** $T_{u,conn} > T_{tra} + T_{inf}$ **then**

8 Add vehicle u to I^t

9 for each selected vehicle $u \in I^t$ **in parallel do**

10 Download the current global model ω^t

11 $\omega_u^{t+1} \leftarrow$ **Vehicle Update** (ω^t, u)

12 Update the global model immediately for each arrived local model ω_u^{t+1}

13 Compute aggregation weight ψ_u^t based on Eq. (12)

14 Update the global model with feature learning based on Eq. (13) - (15)

15 **Return** ω^{t+1}

16 Vehicle Execution:

Input: global model ω^t , regularization parameter λ , decay coefficient δ , $\nabla \zeta_u^t \leftarrow 0$, learning rate η_d

17 Compute the multiplier $\eta_u^t = \eta_d \max\{1, \log(\bar{d}_u^t)\}$

18 Vehicle Update (ω, u):

19 for each local epoch **do**

20 Compute the local loss function based on Eq. (5)

21 Implement regularization based on Eq. (6)

22 Compute $\nabla \zeta_u^t$ based on Eq. (7)

23 Update local model ω_u^t based on Eq. (8)

24 **Terminate** until the maximum number of epochs

25 **Return** ω_u^{t+1}

26 Upload ω_u^{t+1} to the RSU for asynchronous aggregation

corporate weight normalization to reduce computation costs. We use feature extraction on the first layer after the input, and denote the parameters of this layer as $\omega_{(1)}^{t+1}$. For each element $\omega_{(1)}^{t+1}[a, b]$ in column $\omega_{(1)}^{t+1}[b]$ of $\omega_{(1)}^{t+1}$, the updated $\omega_{(1)}^{t+1}$ is obtained by:

$$\kappa_{(1)}^{t+1}[a, b] \leftarrow \frac{\exp\left(\left|\omega_{(1)}^{t+1}[a, b]\right|\right)}{\sum_b \exp\left(\left|\omega_{(1)}^{t+1}[a, b]\right|\right)} \quad (14)$$

$$\omega_{(1)}^{t+1}[a, b] \leftarrow \kappa_{(1)}^{t+1}[a, b] \cdot \omega_{(1)}^{t+1}[a, b] \quad (15)$$

Up to now, the asynchronous FL training during a round is completed. More details of the asynchronous FL are summarized in **Algo. 1**.

B. Hybrid Filter based Popular Content Prediction

Considering the sparseness of content requests, we adopt a hybrid filtering model based on SAE to predict the popular content. The trained SAE model is adopted to extract the

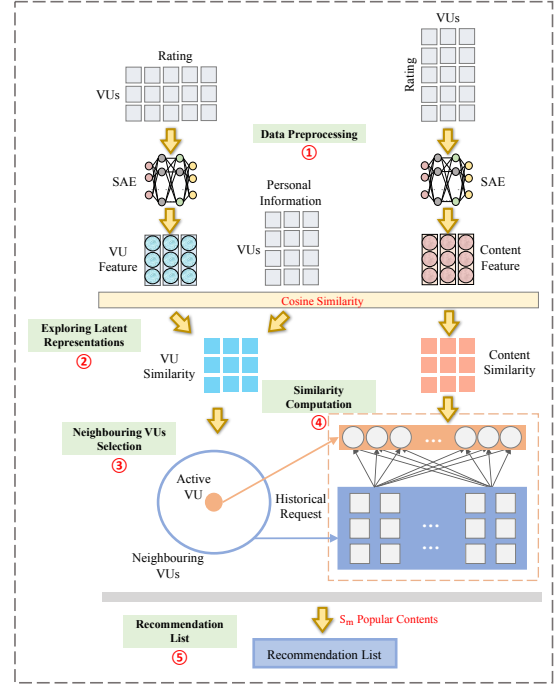


Fig. 3. Popular Content Prediction Process for a Vehicle.

latent feature representation of contents and VUs. Respectively, based on the extracted feature, the potential associations between the content pair and VU pair (*i.e.*, content similarity and VU similarity) can be measured by cosine distance. Through VU similarity, a VU group is formed by randomly selecting an active VU with its k neighbouring VUs. Based on the content similarity of their historical requests, a content recommendation list is generated for RSU to facilitate popular content prediction. The hybrid filtering model combines content-based collaborative filtering and demographic information throughout this process. It relies on similarity metrics to determine the distance between two contents or two VUs based on their contents ratings and personal information. **Fig. 3** illustrates the content prediction process for a vehicle within the cell coverage, consisting of the following procedures:

1) *Data Preprocessing*: Each vehicle generates a rating matrix \mathbb{F} from its local data based on the historical requests of VUs. The rows and columns of this matrix are VUs' ID and their ratings for different contents, respectively. Similarly, a VU information matrix is generated based on the personal information of VUs, where the rows and columns are user IDs and their corresponding personal information, respectively.

2) *Exploring Latent Representations*: Vehicles put the rating matrix as the input for the SAE to explore hidden features among VUs and contents. Along with the user information matrix, these features are used to compute similarity matrices for VUs and contents. Here, we adopt cosine similarity as the metric for similarity due to its effectiveness in sparse matrices. Cosine similarity distinguishes directional differences while being insensitive to absolute value changes, making it well-suited for feature differences measure. Indeed, VU similarity indicates content preference overlap among

VUs, while content similarity represents the level of user interest in different contents.

3) *Neighbouring VUs Selection*: For simplicity, we assume that any current VU is active in this work. Based on the VU similarity matrix, we select the k nearest VUs as the neighbouring VUs of this active VU. Utilizing multiple neighbouring VUs helps balance training errors, thereby improving prediction accuracy. It is worth noting that the growth of k is not unbounded. Further increasing k would be unnecessary once the model converges with sufficient VUs. Next, based on the historical requests, a rating matrix \mathbb{K} is constructed, where the rows and columns represent the IDs of these neighbouring VUs and their ratings for different contents, respectively.

4) *Similarity Computation*: We extract a vector from the rating matrix \mathbb{K} to construct the rating matrix of the active VU. The content similarity matrix allows us to obtain the average similarity between each element in the rating matrix of the active VU and neighbouring VUs.

5) *Recommendation List*: Each vehicle generates a list of the first S_m highest similarity contents. After local training, the list is transmitted to the RSU as a recommendation list. The central RSU aggregates all recommendation lists from vehicles and ranks contents based on their occurrence frequency. The first S_m contents with the highest occurrence frequency are set as predicted popular contents (denoted by set \mathbb{S}).

V. DRL FOR EDGE CACHING DECISION

While the prediction results provide valuable insights for caching decisions, the diversity of content often results in predicted popular content far surpassing the finite cache capacity of an RSU. To efficiently manage the cache space of an RSU, content replacement mechanisms are indispensable in edge caching. Meanwhile, overall cache utilization and effectiveness are prone to underutilization without edge cooperation. Therefore, this section explores intelligent and adaptive caching decisions based on DRL, so as to maximize the benefits of content prediction in IoVs. Notably, to save computing resources, the prediction results generated by the central RSU are directly shared with neighbouring RSUs. From a macroscopic perspective, this makes sense as neighbouring RSUs are in proximity to the central RSU along the same path, and VU preferences are not expected to vary dramatically during a short period.

A. Markov Decision Process

The MDP model is typically used to describe almost all sequential decision-making problems. We model each RSU as an agent for the straightforward definition of state transition. The caching process in an RSU is formulated as an MDP problem. Three critical elements are identified as follows:

State Space: The state is the immediate environment, including relevant information necessary for decision-making. Under the agent setting, the state is represented as $z_i^t = \{a_{\mathcal{F},i,i}^t, q_{i,\mathcal{F}}^t\}$. The former set $a_{\mathcal{F},i,i}^t$ is the current caching state respecting to the contents in RSU i . The latter set $q_{i,\mathcal{F}}^t$ is content requests from vehicles. Specifically, $q_{i,f}^t = 1$ means whether at least one vehicle in RSU i requests for content f .

Action Space: After receiving a state in each slot, the agent is required to make a caching decision. This process begins by sorting the caching contents in RSU i in descending order of popularity, relying on the prediction results provided by the central RSU. With the current state z_i^t , the action d_i^t involves two parts, $a_{\mathcal{F},i,j}^t$ and c_i^t , where matrix $a_{\mathcal{F},i,j}^t$ encodes the decision of RSU i for the requested contents, and decision variable c_i^t indicates the current content replacement control in RSU i . Specifically, if $c_i^t = 1$, the RSU will randomly select $\frac{1}{3}G_i$ contents from those in prediction results \mathbb{S}_i but not currently cached at the edge. These selected contents will replace the least popular contents in its local cache.

Reward Function: The reward value serves as a metric for evaluating the action quality. Our objective is to minimize content delivery delay, in contrast to the goal of maximizing cumulative discounted reward for agents. To align the reward function with our optimization objective, we define the reward function as $r(z_i^t, d_i^t) = e^{-\sum_f \sum_i \sum_j^{N+1} q_{i,f}^t a_{f,i,j}^t \mathcal{T}_{f,i,j}^t}$, where a negative exponential function is adopted to transform the objective legitimately. Notably, punitive negatives will be incorporated into the reward for violating constraint conditions.

B. Value Function Approximation

In general, MDP problems can be tackled through linear or dynamic programming methods. However, these solutions may not always be applicable in dynamic edge caching settings due to their quasi-static and myopic models, which yields unsatisfactory performance [28].

Therefore, we perform the decision through a function approximation structure, which estimates the anticipated future value of executing an action. Specifically, the approximate structure focuses on the agent. Each agent confronts the control and sequential decision-making under uncertainty. At time slot t , the agent observes the current state of the environment as z_t , selects and takes action d^t from the admissible action space based on its policy π , where the policy π is derived as a mapping from current state to corresponding action. After that, the agent transfers to a new state z_{t+1} with the engineered transition function $p(z_{t+1} | z_t, d_t)$. Finally, the agent will receive an immediate reward $r_t = r(z_t, d_t)$ and consider an anticipated reward as it proceeds.

Considering the long-term influences, the recursive state value function $V^\pi(z_t)$ is defined under the policy $\pi(z_t)$. Specifically, it maps state z_t to the expectation of cumulative discounted reward value $U_t = \sum_{t=0}^{\Gamma} \gamma^t r_t$, where discount factor $\gamma \in (0, 1)$ indicates the importance of the predicted future rewards. By applying the Bellman Equation, the state value function is converted to the temporal difference form:

$$\begin{aligned} V^\pi(z_t) &= \mathbb{E}_\pi \left[\left(r(z_t, d_t) + \sum_{t=1}^{\Gamma} \gamma^t r_t \right) \mid z = z_{t+1} \right] \\ &= \left[r(z_t, d_t) + \gamma \sum_{z_{t+1}} p(z_{t+1} | z_t, d_t) V^{\pi^*}(z_{t+1}) \mid z = z_t \right], \end{aligned} \quad (16)$$

where \mathbb{E} represents the expectation.

Along with the process above, the goal of the agent is to make an optimal policy $\pi_i^* \rightarrow d_i^*$ that maximizes the expectation

of cumulative discounted reward value. Consequently, the optimization problem is converted to an optimal value function $V^{\pi^*}(z_t)$ under constraints, which is expressed as:

$$\begin{aligned}
 V^{\pi^*}(z_t) &= \max_{\pi \rightarrow d_t} \left[r(z_t, d_t) + \gamma \sum_{z_{t+1}} p(z_{t+1}|z_t, d_t) V^{\pi^*}(z_{t+1}) \right] \\
 \text{s. t. } &C1 : a_{f,i,j}^t \in \{0, 1\}, \quad \forall f, \forall i, \forall j \\
 &C2 : a_{f,i,j}^t \leq a_{f,j,j}^t, \quad \forall f, \forall i, \forall j \\
 &C3 : \sum_{j=1}^{N+1} a_{f,i,j}^t = 1, \quad \forall f, \forall i \\
 &C4 : a_{f,N+1,N+1}^t = 1, \quad \forall f \\
 &C5 : \sum_{f=1}^F a_{f,i,i}^t \leq G_i, \quad \forall i.
 \end{aligned} \tag{17}$$

This recursive value function equation resolves the sequential decision problem into a series of shorter, tractable time steps, where the action have to be determined in each step. The above constraints are summarized as follows: C1 ensures the integer nature of the binary decision; C2 and C3 are the constraints on conflicting decisions, *i.e.*, each content request can only be served by an RSU with related content or the cloud server ultimately; C4 guarantees that cloud server has cached all contents, and C5 is the cache usage constraint of each RSU.

Explicitly, the control policy π^* , which satisfies Eq. (17) is guaranteed to become the optimal policy. The optimal action for state z_t is easily obtained by $d_t^* = \operatorname{argmax}_{d_t} V^{\pi^*}(z_t, d_t)$.

C. Bellman Update Process

Next, to optimize the strategy and connect the current state with possible actions, we conduct Bellman updates around the action-value function¹ $Q(z_t, d_t)$. This function enables us to estimate the optimal state value function $V^{\pi^*}(z_t)$ under the current state, and the relationship between them is $V^{\pi^*}(z_t) = \max_{\pi} Q^{\pi}(z_t, d_t)$. Therefore, the expected cumulative reward after taking an action d_t is obtained as:

$$\begin{aligned}
 Q^{\pi}(z_t, d_t) &= r(z_t, d_t) + \gamma \sum_{z_{t+1}} p(z_{t+1}|z_t, d_t) V^{\pi^*}(z_{t+1}) \\
 &= r(z_t, d_t) + \gamma \sum_{z_{t+1}} p(z_{t+1}|z_t, d_t) \max_{\pi} Q^{\pi}(z_{t+1}, d_{t+1}).
 \end{aligned} \tag{18}$$

Hereafter, the agent iteratively approximates the action-value function under the current state and executes an action d_t with the highest Q -value. The core process is value iteration, and the iterative formula for Q -value can be obtained as follows:

$$\begin{aligned}
 Q(z_t, d_t)^{new} &= Q(z_t, d_t) + \beta \left[r(z_t, d_t) + \gamma \max_{d_{t+1}} Q(z_{t+1}, d_{t+1}) \right. \\
 &\quad \left. - Q(z_t, d_t) \right],
 \end{aligned} \tag{19}$$

¹No distinction is made between "action-value function" and " Q -function".

Algorithm 2: DRL for Caching Decision-Making

Input: agent set \mathcal{N} , discount factor γ , learning rate β , replay buffer \mathcal{M} , minibatch size J , exploration factor ϵ

Output: neural network parameters θ

- 1 **Each** RSU $i \in \mathcal{N}$ **do**
- 2 **Initialize** experience replay buffer, main neural network with random weight θ , target neural network with $\hat{\theta} = \theta$.
- 3 **for** each episode **do**
- 4 Set $t=0$, obtain the initial state z_0 by randomly caching
- 5 **for** each slots of episode **do**
- 6 Derive an action d_t based on the ϵ -greedy strategy in the current state z_t
- 7 Execute action d_t , observe reward r_t and the next state z_{t+1}
- 8 Store observed experience tuple $\varsigma = \langle z_t, d_t, r_t, z_{t+1} \rangle$ into the replay buffer
- 9 **if** $t \% \text{updateFrequency} == 0$ **then**
- 10 Randomly extract a minibatch of J experience samples from the replay buffer
- 11 **for** each experience sample $\{\varsigma_j\}_{j=1}^J$ **do**
- 12 Update parameter θ by minimizing $\mathcal{L}(\theta)$ as Eq. (21)
- 13 Update parameter $\hat{\theta}$ periodically by θ
- 14 Let $t \leftarrow t + 1$

where $\beta \in (0, 1)$ is the learning rate parameter. The Q -value can definitely converge to the optimal value $Q^{\pi^*}(z_t, d_t)$ when an appropriate β is designed. Each agent regards other agents as part of the environment, and engages in repeated interaction to optimize the edge caching configuration gradually.

D. DRL based Edge Caching Decision-Making

Generally, RL can optimize the Bellman update process through its self-learning ability. However, RL often faces the curse of dimensionality due to the vast action-state space in dynamic scenarios. Besides, recording Q -values in a large table may lead to lengthy searches and memory issues. Thus, DRL is employed for problems in RL trial-and-error interaction.

Throughout the training process, the input training samples that neural networks expect should be distributed independently for exploitation and exploration. Nevertheless, due to the high correlation of continuous states, action-value function estimations through on-policy updates is easy to perform non-uniform overestimation. This will introduce instability and adverse impacts on the actual value. Therefore, we adopt off-policy updates in practice. The basic principle in our proposed DRL method is to separate the policy of behavior and target by different state-action functions. By adopting neural network-based parameter approximations, the agent expects to make an optimal action by solving the sequential optimization in Eq. (19). It estimates the anticipated value of decisions with

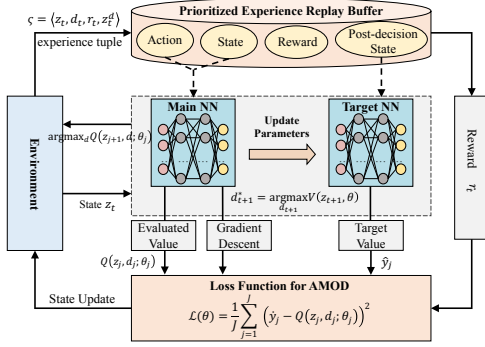


Fig. 4. The Update Process of Action-Value Function.

spatiotemporal dependencies, and approximates the optimal value $Q^*(z_t, d_t)$ with reconstructed Q -function $Q(z_t, d_t; \theta)$. Hereafter, the agent actuates the saved Q -value towards the target value through parameter updates.

Overall, the value update process in our proposed method largely follows the Double Deep Q Network (DDQN). Instead of performing bootstrapping directly, a target network mechanism exists to disrupt the relevance. Specifically, two neural networks with the same structure but different parameters are maintained. The target neural network aims to acquire the temporal difference within the one-step return value, while the main neural network evaluates the current Q -value. For the sake of learning stability, the weight parameters $\hat{\theta}$ of the target neural network are updated periodically (spaced a few training steps) by the counterpart θ of the main neural network. Their update rule follows $\hat{\theta} = \zeta\theta + (1 - \zeta)\hat{\theta}$ with $\zeta \ll 1$. Note that the one-step return method (*i.e.*, TD(0)) in the target neural network is simply based on the immediate reward value. Then, the target network mechanism is used to generate the one-step target Q -value \hat{y}_t as:

$$\begin{aligned} \hat{y}_t &= r_t + \gamma \max_d Q(z_{t+1}, d; \hat{\theta}) \\ &= r_t + \gamma Q[z_{t+1}, \arg\max_d Q(z_{t+1}, d; \theta_t), \hat{\theta}_t]. \end{aligned} \quad (20)$$

Furthermore, to ensure the stability of this neural network-based nonlinear approximation, we utilize a prioritized experience replay buffer \mathcal{M} to reuse the pre-existing experiences and break the correlation among data training. During training sample collection, the agent stores its experience tuple $\varsigma = \langle z_t, d_t, r_t, z_{t+1} \rangle$ into the experience replay buffer, which involves its current state and available action set. The collected experience samples are utilized to train the neural network parameters for value approximation. This entails extracting a minibatch of previous experience samples from the replay buffer at each iteration, enabling the agent to reinforce its knowledge through replaying.

Since $Q(z_t, d_t; \theta_t) = \hat{y}_t$ is possible with a smaller error, we attempt to optimize parameters for the Q -function approximation. The loss function is interpreted as the Euclidean distance between target value and estimated value. For a minibatch of experience sample $\{\varsigma_j\}_{j=1}^J$, it is converted as:




Fig. 5. Content access delay versus episodes.

$$\begin{aligned} \mathcal{L}(\theta) &= \frac{1}{J} \sum_{j=1}^J \left(r_j + \gamma Q \left(z_{j+1}, \arg\max_d Q(z_{j+1}, d; \theta_j), \hat{\theta}_j \right) \right. \\ &\quad \left. - Q(z_j, d_j; \theta_j) \right)^2, \end{aligned} \quad (21)$$

where each experience samples $\varsigma_j = \langle z_j, d_j, r_j, z_{j+1} \rangle$ is used to update the parameter θ toward the target value by minimizing loss function $\mathcal{L}(\theta)$, and a gradient guiding updates of θ can be calculated by $\frac{\partial \mathcal{L}(\theta)}{\partial \theta}$.

Importantly, the optimal action is typically constrained within a finite search space, which depends on the quantity and quality of the training data. To introduce randomness in the action selection, we apply and modify the ϵ -greedy strategy to balance exploitation and exploration, with ϵ as a decreasing parameter. Specifically, the agent gets the action that maximizes the Q -value with the probability $1 - \epsilon$ (exploiting), and the other action with the tiny probability ϵ (exploring).

Fig. 4 shows the update process of Q -value, and more details are summarized in Algo. 2. We estimate action-value function for each time slot under the parameter approximation. By the constructed loss function, the neural network updates parameters and further obtains optimal decisions. Finally, to continue the process to the next slot, the aforementioned value function approximation is solved again until a finite time horizon within an episode. Value iteration is terminated when it reaches the given maximum number of episodes, and the desired action-value function becomes stable as well.

VI. PERFORMANCE EVALUATION

We consider a one-way highway scenario, with all experiments conducted in Python 3.8 and tested on a processor (AMD® Threadripper™ PRO5995WX@2.70GHz). The system comprises a cloud server and 5 RSUs, where the cell coverage is 500 m. Unless otherwise stated, the vehicle density is 15 vehicles/km, and the initial cache capability of each RSU is 100 contents. Table II provides details of other parameters.

The MovieLens 1M dataset [29] is utilized to emulate the content request of VUs. The MovieLens 1M dataset comprises 1,000,209 ratings from 6040 anonymized users for 3952 movies. The ratings range from 0 to 5. Each dataset entry includes indices for both user and movie, alongside the timestamp of the rating. Furthermore, the dataset encompasses

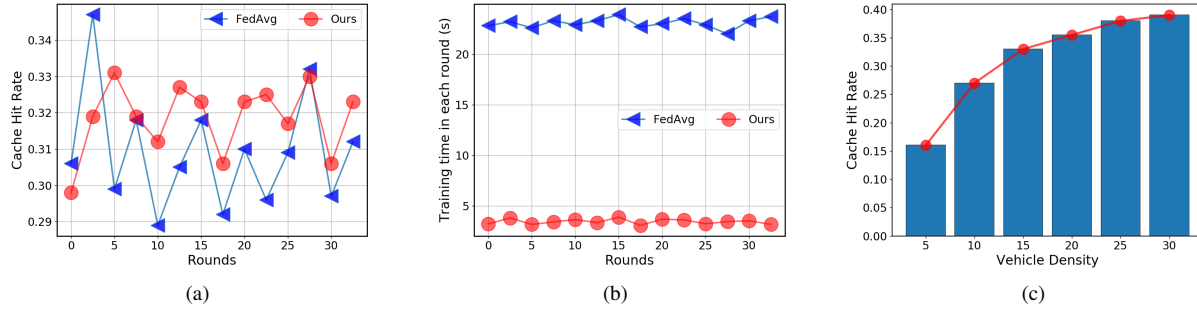


Fig. 6. (a) Content access delay versus episodes. (b) Cache hit rate versus the round in FL. (c) The training time for each round in FL. (d) Cache hit rate versus the vehicle density.

TABLE II
PARAMETERS SETUP

Parameter	Value	Parameter	Value
System Parameters			
V_{\max}, V_{\min}	60km/h, 100km/h	s_f	0.5MB
σ^2	-95dBm	B_i	2MHz
$T_{i,N+1}$	80ms	$T_{i,j}$	15ms
P_i	38dBm		
FL Parameters			
η_d	0.001	T_{tra}	3s
δ	0.001	T_{inf}	0.5s
α	0.4	k	8
λ	0.005		
DRL Parameters			
M	500	J	32
β	3×10^{-4}	ϵ	0.05
γ	0.98		

pertinent user demographic information (gender, age, occupation, *etc.*). Particularly, we assume that a rating for a movie equates to a content request for that particular content. This conjecture aligns with the rationality that users generally rate movies they have already seen.

For performance comparison, the following baselines are introduced:

- 1) **Oracle [30]**: Oracle is an omniscient algorithm that the knowledge of future demands is known in advance. It yields a performance upper bound for other baselines.
- 2) **Least Recently Used (LRU)**: LRU replaces the cached content with the longest un-accessed time firstly in RSU.
- 3) **NoDRL**: NoDRL only predicts popular content but also doesn't use DRL for caching decisions. RSUs cache contents randomly from the predicted popular ones.
- 4) **Federated Averaging (FedAvg) [31]**: FedAvg is an synchronous FL method, where the RSU use a weighted average to update the global model.

Quantitative metrics: 1) Cache hit rate: the ratio of requests served by an RSU to the total number of requests. 2) Content access delay: the average delay to access a content.

Fig. 5 presents the convergence curve, indicating a gradual decrease in content access delay, stabilising after about 150 episodes. This shows the learning process of RSU, which can achieve the optimal caching strategy after 150 episodes.

Fig. 6(a) shows the cache hit rate comparison between our proposed method and FedAvg versus rounds. Our proposed method maintains a relatively stable cache hit rate within a

small fluctuation range (29.7% to 33.2%) during the 30 rounds. FedAvg exhibits up to 6% fluctuation in cache hit rate over the same period. This indicates that our proposed method slightly outperforms FedAvg regarding model prediction accuracy. This improvement can be attributed not only to the accuracy improvement through the mobility-aware asynchronous aggregation in our method, but also to a series of measures we adopted to cope with the learning instability caused by stragglers. Similarly, **Fig. 6(b)** compares the training time for each round between our proposed method and FedAvg. It is evident that our proposed method significantly reduces the training time due to asynchronous aggregation. The training time for FedAvg ranges from 22.1 to 23.9 seconds, as it must wait for all local updates from vehicles before aggregation. Besides, **Fig. 6(c)** illustrates that vehicle density positively impacts the cache hit rate. This is reasonable because as vehicle density increases, the amount of training data and uploaded local updates in the cell coverage also increase, further enhancing the model prediction accuracy.




Fig. 7. (a) Cache hit rate versus the cache capacity of RSUs. (b) Content access delay versus the cache capacity of RSUs.

Next, we study the effect of varying cache capacities on different methods, as shown in **Fig. 7(a)** and **7(b)**. Increasing RSU capacity resulted in higher cache hit rates and lower access delay across all methods. As expected, Oracle achieves the best performance, but its practicality is limited. Among other methods, our proposed method consistently outperforms them in terms of quantitative metrics. For example, when the cache capacity reaches 350 MB, the edge hit rate is improved by roughly 6%, 21%, and 15% compared to FedAvg, LRU, and NoDRL, respectively. At this moment, content replacement processes may infrequently occur in the RSUs.

VII. CONCLUSION

This paper investigated a mobility-aware edge caching strategy by exploiting asynchronous FL and DRL. Firstly, popular content prediction was conducted using the SAE model. We employed an asynchronous FL framework for local updates and global aggregation of SAE models. The trained SAE model extracts latent features from VUs and contents. These features construct a hybrid filtering model for popular content recommendation. Secondly, to maximize the benefits of content prediction, we explored intelligent caching decisions considering content delivery and cache replacement. Based on the formulated MDP problem, we proposed a DRL-based edge caching decision-making and adopted neural network-based parameter approximations. Extensive simulations were conducted based on real-world data trajectory.

REFERENCES

- [1] P. Arthurs, et al., "A Taxonomy and Survey of Edge Cloud Computing for Intelligent Transportation Systems and Connected Vehicles," *IEEE Trans. Intell. Trans. Syst.*, vol. 23, no. 7, pp. 6206–6221, 2022.
- [2] J. Zhang, et al., "Mobile Edge Intelligence and Computing for the Internet of Vehicles," *Proc. IEEE*, vol. 108, no. 2, pp. 246–261, 2020.
- [3] P. Wang, et al., "Marginal Value-Based Edge Resource Pricing and Allocation for Deadline-Sensitive Tasks," *Proc. IEEE Int. Conf. Comput. Commun. (INFOCOM)*, pp. 1–10, 2023.
- [4] L. Su, et al., "Data and Channel-Adaptive Sensor Scheduling for Federated Edge Learning via Over-the-Air Gradient Aggregation," *IEEE Internet of Things J.*, vol. 9, no. 3, pp. 1640–1654, Feb. 2022.
- [5] K. Jiang, et al., "Intelligence-empowered mobile edge computing: Framework, issues, implementation, and outlook," *IEEE Network*, vol. 35, no. 5, pp. 74–82, 2021.
- [6] J. Zhao, et al., "Edge Caching and Computation Management for Real-Time Internet of Vehicles: An Online and Distributed Approach," *IEEE Trans. Intell. Trans. Syst.*, vol. 16, no. 10, pp. 2183–2197, 2021.
- [7] H. Zhou, et al., "Incentive-driven Deep Reinforcement Learning for Content Caching and D2D Offloading," *IEEE J. Sel. Areas Commun.*, vol. 39, no. 8, pp. 2445–2460, 2021.
- [8] B. Jedari, G. Premsankar, G. Illahi, M. D. Francesco, A. Mehrabi, and A. Ylä-Jääski, "Video Caching, Analytics, and Delivery at the Wireless Edge: A Survey and Future Directions," *IEEE Commun. Surveys & Tutorials*, vol. 23, no. 1, pp. 431–471, 2021.
- [9] D. Wang, et al., "Secure Long-Range Autonomous Valet Parking: A Reservation Scheme With Three-Factor Authentication and Key Agreement," *IEEE Trans. Veh. Technol.*, vol. 72, no. 3, pp. 3832–3847, 2023.
- [10] H. Li, K. Ota, and M. Dong, "Deep Reinforcement Scheduling for Mobile Crowdsensing in Fog Computing," *ACM Trans. Internet Techn.*, vol. 19, no. 2, pp. 1–18, 2019.
- [11] Z. Yu, et al., "Mobility-aware proactive edge caching for connected vehicles using federated learning," *IEEE Trans. Intell. Trans. Syst.*, vol. 22, no. 8, pp. 5341–5351, 2021.
- [12] D. Qiao, et al., "Adaptive Federated Deep Reinforcement Learning for Proactive Content Caching in Edge Computing," *IEEE Trans. Parallel. Distrib. Syst.*, vol. 33, no. 12, pp. 4767–4782, 2022.
- [13] Y. Chen, et al., "Asynchronous online federated learning for edge devices with non-IID data," in *Proc. IEEE Int. Conf. Big Data*, pp. 15–24, 2020.
- [14] H. Zhou, et al., "Distributed Multi-Agent Reinforcement Learning for Cooperative Edge Caching in Internet-of-Vehicles," *IEEE Trans. Wire. Commun.*, early access, 2023, doi: 10.1109/TWC.2023.3272348.
- [15] L. T. Tan, and R. Q. Hu, "Mobility-aware edge caching and computing in vehicle networks: A deep reinforcement learning," *IEEE Trans. Veh. Technol.*, vol. 67, no. 11, pp. 10190–10203, Jul. 2018.
- [16] H. Zhang, et al., "Power Control Based on Deep Reinforcement Learning for Spectrum Sharing," *IEEE Trans. Wireless Commun.*, vol. 19, no. 6, pp. 4209–4219, 2020.
- [17] J. Shu, et al., "Clustered Federated Multitask Learning on Non-IID Data With Enhanced Privacy," *IEEE Internet of Things J.*, vol. 10, no. 4, pp. 3453–3467, 2023.
- [18] V. Smith, et al., "Cocoa: A general framework for communication-efficient distributed optimization," *Journal of Machine Learning Research*, vol. 18, no. 230, pp. 1–49, 2018.
- [19] X. Jin, et al., "Collaborating between local and global learning for distributed online multiple tasks," in *Proc. ACM Int. Conf. Inf. Knowl. Manage. (CIKM)*, pp. 113–122, 2015.
- [20] N. Aybat, et al., "An asynchronous distributed proximal gradient method for composite convex optimization," in *Proc. Int. Conf. Mach. Learn. (ICML)*, pp. 2454–2462, 2015.
- [21] C. Xie, et al., "Asynchronous Federated Optimization," in *Proc. NeurIPS Workshop on Optimization for Machine Learning (OPT)*, pp. 1–11, 2020.
- [22] L. Zhao, et al., "Towards cooperative caching for vehicular networks with multi-level federated reinforcement learning," in *Proc. IEEE Int. Conf. Commun. (ICC)*, pp. 1–6, 2021.
- [23] K. Wang, et al., "An Efficient Content Popularity Prediction of Privacy Preserving Based on Federated Learning and Wasserstein GAN," *IEEE Internet of Things J.*, vol. 10, no. 5, pp. 3786–3798, 2023.
- [24] Y. Jiang, et al., "Federated Learning-Based Content Popularity Prediction in Fog Radio Access Networks," *IEEE Trans. Wireless Commun.*, vol. 21, no. 6, pp. 3836–3849, 2022.
- [25] G. Qiao, et al., "Deep reinforcement learning for cooperative content caching in vehicular edge computing and networks," *IEEE Internet of Things J.*, vol. 7, no. 1, pp. 247–257, 2020.
- [26] L. Chen, et al., "A3C-Based and Dependency-Aware Computation Offloading and Service Caching in Digital Twin Edge Networks," *IEEE Access*, vol. 11, no. 10, pp. 57564–57573, 2023.
- [27] V. Kirilin, et al., "RLcache: Learning-based cache admission for content delivery," *IEEE J. Sel. Areas Commun.*, vol. 38, no. 10, pp. 2372–2385, 2020.
- [28] N. Luong, et al., "Applications of deep reinforcement learning in communications and networking: A survey," *IEEE Commun. Surveys Tuts.*, vol. 21, no. 4, pp. 3133–3174, May 2019.
- [29] F. M. Harper, and J. A. Konstan, "The movielens datasets: History and context," *ACM Trans. Interact. Intell. Syst.*, vol. 5, no. 4, pp. 1–19, 2016.
- [30] X. Wang, et al., "Federated Deep Reinforcement Learning for Internet of Things With Decentralized Cooperative Edge Caching," *IEEE Internet of Things J.*, vol. 7, no. 10, pp. 9441–9455, 2020.
- [31] J. Konecny, et al., "Federated optimization: Distributed machine learning for on-device intelligence," *arXiv:1610.02527. [Online]*, 2016. Available: <https://arxiv.org/abs/1610.02527>.